

How Far Can a PSRAM-less ESP32 Go? A TinyML Deployment Study for WiFi CSI Human Activity Recognition

ANNAS WASEEM MALIK¹ and MUHAMMAD AHMAD¹

¹Faculty of Information Technology, University of Central Punjab, Lahore, Pakistan (e-mail: annas.waseem@ucp.edu.pk; muhammad.ahmad@ucp.edu.pk)

Corresponding author: Annas Waseem Malik (e-mail: annas.waseem@ucp.edu.pk).

● **ABSTRACT** WiFi Channel State Information (CSI) is a useful signal for device-free human activity recognition (HAR), and a growing number of studies move the inference onto the cheap microcontroller-class radios themselves so the technology can be deployed at scale. Almost all of that on-device work, however, targets the more capable ESP32-S3 with external PSRAM, or it uses the microcontroller only as a CSI collector and runs the model elsewhere. The cheapest, most common, and weakest member of the family, the classic ESP32 with no PSRAM and only about 150 kilobytes of usable contiguous internal SRAM, has not been characterized as an inference target for CSI HAR. This paper provides that characterization and reaches a two-part conclusion: on this device, getting a model to run is not the bottleneck, but recognizing an unseen person is. For two public datasets (UT-HAR and CSI-HAR) we train and quantize a graded set of models, from classical learners (Decision Tree, Random Forest, and a small multilayer perceptron on handcrafted statistics) through a tiny-CNN width sweep up to five standard deep architectures (convolutional, recurrent, attention, Kolmogorov-Arnold, and state-space), reporting accuracy, model complexity, 8-bit conversion feasibility, and on-device fit. On the deployment side the board is more capable than assumed. A convertibility wall blocks only the recurrent and state-space families, which do not convert to the TensorFlow Lite for Microcontrollers builtin operator set on either dataset. By flashing every convertible model to the board we find that a full deep convolutional network and a convolution-augmented Transformer both run on the bare classic ESP32 within 10 to 54 kilobytes of tensor arena, alongside the classical learners and the tiny CNNs (int8 models of 9 to 38 kilobytes, 14 to 117 milliseconds per inference). There is no general memory wall, which corrects a common assumption. The binding constraint is accuracy on a new person. Under a subject-independent leave-one-user-out protocol on CSI-HAR, the same models that score 69 to 96% on a random split, which leaks each user across the train and test sets, fall to 41 to 63%, a mean drop of 29 percentage points. The contribution is a fair, reproducible deployment map of the most constrained commodity WiFi microcontroller, and the evidence that its open problem is cross-subject generalization rather than fitting the model on the chip, with all code, models, and measurement firmware released.

● **INDEX TERMS** Channel state information, decision tree, edge computing, ESP32, human activity recognition, int8 quantization, Kolmogorov-Arnold networks, microcontrollers, on-device inference, random forest, state-space models, TinyML, WiFi sensing.

I. INTRODUCTION

WiFi Channel State Information (CSI) has become a practical signal for device-free human activity recognition (HAR). Because CSI is reflected and scattered by a moving body, a receiver can infer activity without a camera and without any wearable on the user, which is attractive for privacy-sensitive settings such as homes, clinics, and

assisted-living facilities. The release of commodity CSI extraction tools [1] and a decade of deep-learning architectures [2], [3] have pushed in-distribution accuracy on the standard benchmarks above 98%.

A second trend is just as important for real deployment: moving the inference onto the cheap radios themselves. The Espressif ESP32 family, dual-core microcontrollers with

WiFi and a few hundred kilobytes of RAM that cost under ten US dollars, is a widespread platform for this. Running the model on the device avoids streaming raw CSI to a server, removes the associated privacy and bandwidth costs, and is what would make WiFi sensing deployable at the scale of a light switch. Recent work has shown that compact convolutional and recurrent models can be quantized to 8-bit precision and run on ESP32-class hardware [15], [17].

These two trends meet at a question that the literature has not answered. The published on-device CSI studies target the ESP32-S3, which carries several megabytes of external PSRAM and therefore relaxes the memory constraint, or they use the microcontroller purely as a CSI collector and run the classifier on a gateway or a server. The classic ESP32, the original and still the most widely deployed member of the family, has no PSRAM and exposes only a few hundred kilobytes of internal SRAM, of which a single contiguous block of roughly 150 kilobytes is realistically available to a model at runtime once the radio stacks are disabled. It is also, by Espressif's own ranking, the weakest CSI source in the family. This cheapest and most constrained device is precisely the one a cost-sensitive deployment would choose, and it has not been characterized as an inference target for CSI HAR. A practitioner who holds a bare classic ESP32 has no evidence about what, if anything, will actually run on it.

A. CONTRIBUTIONS

This paper answers that question with a deployment-characterization study rather than a new model. It reaches a two-part result: the bare classic ESP32 runs far more than expected, so the limiting factor for practical CSI HAR on this hardware is not deployment but cross-subject generalization. The contributions are:

- 1) A **deployability frontier** for WiFi CSI HAR on the PSRAM-less classic ESP32, spanning classical learners, tiny neural networks, and five standard deep architectures, each reported with accuracy, complexity, 8-bit conversion feasibility, model size, and on-device fit.
- 2) An **on-device characterization of which deep architectures actually deploy**, measured by flashing each to the board: the convolutional and Transformer models convert and run on the bare device, the recurrent and state-space models do not convert (a convertibility wall), and the Kolmogorov-Arnold model converts but fails to allocate at runtime. This corrects the common assumption that deep models cannot fit such a device.
- 3) A **positive deployability result**: classical learners and tiny convolutional networks both fit the device, the former trivially and the latter as int8 models of 9 to 38 kilobytes, and we quantify the accuracy each retains on two datasets.
- 4) Identification of the **cross-subject generalization gap** as the binding constraint on this hardware. Under a subject-independent leave-one-user-out protocol on CSI-HAR, six models drop by a mean of 29 percentage points (the best from 96% to 63%) relative to a

random split that leaks each user across train and test. Deployment is solved on this device; generalization to an unseen person is not.

- 5) A **reproducible artifact set**: the training and quantization notebook, the exported int8 models and classical models, the emlearn C export of the trees, and the ESP32 measurement firmware and protocol.

B. SCOPE AND HONEST POSITIONING

We do not claim a new model or a new accuracy record; UT-HAR accuracy is saturated and is not where the contribution lies. The closest prior work is the on-device CNN and DenseNet study of Hernandez et al. [15] on the PSRAM-equipped ESP32-S3, the STAR framework [16], which runs a lightweight CSI model on an edge NPU and uses an ESP32-S3 only to collect CSI, and the ESP-Fi HAR dataset and benchmark [17] on the ESP32-C3. None of these characterizes inference on the bare, PSRAM-less classic ESP32, and they cover convolutional and recurrent models only. The closest TinyML work in spirit, Suwannaphong et al. [19], optimizes models for 32 to 64 kilobytes of RAM but for indoor localization, not CSI HAR. Our contribution is to characterize the bare classic ESP32, the most constrained commodity target, across a graded model set and to report the deployment walls and the positive frontier that a GPU-only or PSRAM-equipped study does not reveal. We regard a fair, reproducible deployment map, including the negative findings, as a useful contribution for practitioners.

II. RELATED WORK

A. DEEP LEARNING FOR WIFI CSI HAR

The move from hand-crafted features to end-to-end learning produced a sequence of architectures evaluated on a stable set of benchmarks. ABLSTM [4] introduced attention-augmented bidirectional recurrence; THAT [5] replaced recurrence with a two-stream convolution-augmented Transformer; the SenseFi library [3] standardized training and evaluation across UT-HAR, NTU-Fi, and Widar and reports strong baselines (ResNet-18 at 98.11% on UT-HAR). Recent work emphasizes lightweight and complementary-feature models, including a pruned attention-GRU [6] and the amplitude-and-phase PA-CSI model [7]. Almost all of this work optimizes accuracy on server-class hardware, on which the binding constraint of this paper, on-chip memory, does not apply.

B. KOLMOGOROV-ARNOLD NETWORKS AND STATE-SPACE MODELS

Kolmogorov-Arnold Networks [8] replace the fixed activations of a multilayer perceptron with learnable univariate functions on the network edges, typically parameterized by splines or orthogonal polynomials. They have been applied to sensor-based HAR, for example on accelerometer data [10], but their use for WiFi CSI, which we include in our deep tier, is still nascent and has not been measured on

a microcontroller. On the gateway side, recent consumer-healthcare WiFi HAR favours advanced architectures such as neurosymbolic CNN-Transformer models [9] rather than microcontroller-deployable ones. State-space models such as Mamba [11] model long sequences with a linear-time recurrence and have been applied to CSI HAR on the GPU by TF-Mamba [12] and bidirectional-Mamba HAR [13]. Mamba has been ported to the ESP32-S3 for keyword spotting and inertial HAR by MambaLite-Micro [14], but not for WiFi CSI, and its selective scan is known to be awkward for the microcontroller runtime. Whether either family converts to and fits the bare classic ESP32 for CSI HAR has not been measured.

C. CLASSICAL AND TINYML MODELS ON MICROCONTROLLERS

Outside the deep-learning literature, classical learners remain the workhorses of microcontroller deployment because their footprint and inference cost are small and predictable. Decision trees and tree ensembles compile to compact branch-only C code through tools such as emlearn [21], and prior CSI work has shown that classical models on statistical features reach useful accuracy off-device [22]. TinyML neural-architecture-search studies such as TinyTNAS [23] and MicroNAS [24] target microcontroller HAR, but for inertial sensors rather than CSI, and Suwannaphong et al. [19] optimize TinyML models for 32 to 64 kilobytes of RAM for indoor localization. The present study places classical learners, tiny networks, and deep networks on the same frontier for CSI HAR on the most constrained ESP32.

D. ON-DEVICE WIFI SENSING AND QUANTIZATION

The reference on-device study is Hernandez et al. [15], who deploy a quantized LSTM-DenseNet on the ESP32-S3 at 92.43% accuracy, 232 ms latency, and 127 kB memory using TensorFlow Lite for Microcontrollers and 8-bit quantization, on a board with external PSRAM. STAR [16] runs a lightweight CSI HAR model on an edge NPU with hardware-aware co-optimization, using an ESP32-S3 only as the CSI front end, and the ESP-Fi HAR dataset [17] provides an ESP32-C3 corpus with benchmark deep models. A recent survey catalogues energy-efficient WiFi sensing [18]. Across this line, 8-bit integer quantization is effectively mandatory for the ESP32 memory budget, and the question of which models survive it on the PSRAM-less device is open. The non-convertibility of recurrent and dynamic operators to the microcontroller runtime is documented in the TensorFlow and TensorFlow Lite Micro issue trackers [26], [27], which we corroborate here for CSI models.

Gap. No prior study characterizes the bare, PSRAM-less classic ESP32 as an inference target for WiFi CSI HAR, nor maps which models across the classical, tiny, and deep spectrum actually fit it. This paper fills that gap.

III. MODELS UNDER STUDY

The frontier spans three tiers (Figure 1), all consuming the same input representation. A CSI window is a real-valued tensor $\mathbf{X} \in \mathbb{R}^{T \times F}$ with T time steps and F per-step features. We downsample the time axis to $T = 64$ by uniform index selection, which costs almost no accuracy on UT-HAR while shrinking activations, and we standardize each window.

A. CLASSICAL LEARNERS

For the classical tier, each window is summarized by five statistics computed over time for every feature dimension: mean, standard deviation, minimum, maximum, and range, giving a compact $F \times 5$ vector that is cheap to compute on a microcontroller. On these features we train a depth-limited Decision Tree, a Random Forest of twenty depth-limited trees, and a small multilayer perceptron with one hidden layer of thirty-two units. Trees and forests are exported to branch-only C with emlearn [21], which reproduces the trained decision structure; we therefore report the off-device accuracy for these models, and confirming that emlearn's integer feature handling leaves accuracy unchanged on the device is a minor remaining check.

B. TINY NEURAL NETWORKS

The tiny tier comprises a tiny multilayer perceptron (one hidden layer of thirty-two units on the flattened window) and a tiny-CNN width sweep with 8, 16, and 32 channels in the first convolutional block. Each tiny CNN has two 1D convolutional blocks, global average pooling, and a dense classifier. These models are trained in floating point and then quantized to int8.

C. DEEP ARCHITECTURES

The deep tier comprises five standard families: a 1D convolutional network, a bidirectional gated recurrent unit, a convolution-augmented Transformer following the spirit of THAT [5], a Chebyshev Kolmogorov-Arnold network [8], and a lightweight diagonal state-space model [11]. They are sized to comparable parameter budgets and contribute the upper, large-model part of the frontier. To keep the comparison fair and the results verifiable, the deep tier uses the same training, 8-bit quantization, and on-device measurement protocol as the classical and tiny tiers (Section IV); the deep-tier training notebook and the resulting int8 models are part of the released artifact set, so every value in Table 2 is reproducible from the release rather than taken on trust.

IV. EXPERIMENTAL SETUP

A. DATASETS

We use two public datasets so that the frontier is not tied to a single corpus. UT-HAR, obtained through the SenseFi loaders [3], is the most widely used public CSI HAR benchmark (Intel 5300, seven activities, native $T = 250$, $F = 90$). CSI-HAR [25] is an ESP32-collected dataset of seven activities (bend, fall, lie down, run, sit down, stand up, walk) recorded by three users with twenty samples each. Every recording is a




FIGURE 1. System and measurement pipeline. Models are trained and quantized off-device on a cloud GPU, then exported as int8 TensorFlow Lite (neural models) or emlearn C (classical models), deployed on the bare classic ESP32, and measured for on-device latency and peak RAM. Convertibility and on-device allocation are evaluated at the deploy stage, and the result is the deployability frontier.




FIGURE 2. The ESP32-WROOM-32 development board (ESP32-D0WD-V3 SoC, dual-core Xtensa LX6, 520 kB SRAM, no PSRAM, CP2102 USB-to-UART bridge) used for all on-device measurements. The board is the bare classic ESP32 and is connected over USB for flashing and serial readout.

variable-length sequence over 52 subcarriers. UT-HAR uses its standard split, and each model is trained over three seeds. CSI-HAR is evaluated subject-independently with leave-one-user-out: each of the three folds trains on two users and tests on the held-out third, and we report the mean and standard deviation across folds. This avoids the subject leakage of a random within-user split and is the standard rigour for HAR generalization. Both datasets are z-score normalized per window and resampled to $T = 64$.

B. TRAINING AND QUANTIZATION

All neural models are trained with one shared pipeline: Adam at learning rate 10^{-3} , batch size 64, for 40 epochs, with sparse categorical cross-entropy on logits. Classical learners are trained over three seeds and we report mean and standard deviation. Each trained network is converted to 8-bit integer precision with TensorFlow Lite using post-training static quantization calibrated on 300 representative training windows; we attempt full int8 input and output first and fall back if an operator forbids it. We record, for each model, whether conversion succeeds and the resulting model size. Off-device training and quantization run on a cloud GPU; the released notebook reproduces every off-device number.

C. ON-DEVICE PLATFORM AND MEASUREMENT

The deployment target is a commodity classic ESP32 on an ESP32-WROOM-32 module (Figure 2; ESP32-D0WD-V3 SoC, dual-core Xtensa LX6 at 240 MHz, 520 kB internal SRAM, no PSRAM, 4 MB flash) running TensorFlow Lite for Microcontrollers. With the WiFi and Bluetooth stacks disabled and the interpreter arena allocated from the largest free internal block, a single contiguous block of roughly 150 kilobytes is available to a model at runtime (the firmware reports a largest free internal block of about 152 kB); this is the binding budget. For each model that converts we attempt allocation on the device. When allocation succeeds, we measure inference latency averaged over 1000 invocations using the hardware timer, peak RAM from the interpreter's reported arena usage, and flash from the model byte array. The board carries no current-sense hardware, so rather than measure energy we estimate it as the measured latency times the ESP32 active power. The ESP32 datasheet [20] lists 30 to 68 mA in modem-sleep with the CPU at 240 MHz and the radio off, which is the condition during our inference; we take a representative 50 mA at 3.3 V (about 0.165 W). The resulting per-inference energy in Table 3 is therefore an estimate, not a direct measurement, and a current-probe study is left to future work. Decision trees and forests are deployed as emlearn C and timed the same way. The firmware, flash offsets, and capture protocol are released with the paper.

D. METRICS

Accuracy measures recognition quality. For classical learners we report the mean and standard deviation over three seeds. Model complexity is reported as tree node count, total forest node count, or parameter count, the quantity that maps to footprint. The int8 model size and the conversion outcome measure deployability off-device, and latency, RAM, flash, and on-device fit measure it on the device. The accuracy-versus-cost frontier is the main result.

V. RESULTS

We first report the classical and tiny tiers, then the deep tier and which of its families actually deploy on the device, and

finally the on-device latency and memory measurements.

A. THE DEPLOYABLE FRONTIER: CLASSICAL AND TINY MODELS

Table 1 reports the classical and tiny tiers on both datasets. Every model in this table is small enough to be a realistic candidate for the device. The two datasets sit in different accuracy regimes. UT-HAR is evaluated on its fixed split and is saturated: the 32-channel tiny CNN reaches 96.7%, the 16-channel tiny CNN 94.9%, and even the classical multilayer perceptron on statistics reaches 95.3%, so the small models are within a point or two of the deep models on this benchmark. CSI-HAR is evaluated subject-independently (leave-one-user-out), and accuracy is much lower for every model, from 41.4% for the Decision Tree to 62.1% for the 32-channel tiny CNN. This gap is not a deployment artefact: it reflects how hard cross-subject generalization is on a small three-user CSI corpus, and it is the honest number a deployed device would face on an unseen person. The standard deviations confirm it: CSI-HAR variance across held-out users is large (up to ± 6.6 points), whereas UT-HAR variance across seeds is small. Within each dataset the cost ordering is consistent: wider tiny CNNs are more accurate, the 32-channel model is best on both datasets, and the tiny multilayer perceptron is the weakest network for its size because flattening the window inflates its parameter count without a matching accuracy gain. Figure 4 plots accuracy against int8 size and shows the two regimes and the within-dataset width ordering.

Figure 5 compares accuracy across the classical and tiny tiers and the two datasets, and Figure 6 shows model complexity on a logarithmic axis, with the classical learners one to three orders of magnitude below the networks. Accuracies are lower on CSI-HAR than on UT-HAR throughout, which is expected: CSI-HAR is collected from a single classic ESP32, the weakest CSI source, has only three users and few samples per class, and is evaluated subject-independently rather than on a random split. It is therefore the more honest stress test of what a deployed device will face on an unseen person, and a random within-subject split (which we deliberately avoid) would overstate accuracy by leaking each user across train and test, an effect we quantify in Section V-D.

B. THE DEEP TIER: WHICH FAMILIES DEPLOY

Table 2 reports the five deep architectures on both datasets together with the result of flashing every convertible model to the physical device. The outcome is more nuanced than a blanket “deep models do not fit,” and we report it as we measured it.

One barrier is hard and architectural, a convertibility wall. The bidirectional GRU and the lightweight state-space model, despite competitive floating-point accuracy, do not convert to a microcontroller-deployable model on either dataset. The recurrent model compiles to the fused CudnnRNNV3 operator and the state-space model to dynamic TensorList operators, neither of which is in the

TensorFlow Lite for Microcontrollers builtin operator set; both require the Select-TF-ops path that the ESP32 runtime does not provide [26], [27]. This wall is intrinsic to the architecture, so it holds on both datasets.

The feed-forward families, in contrast, mostly do deploy, which corrects an assumption we held earlier. The convolutional network allocates and runs within a 10 kB tensor arena on UT-HAR and 12 kB on CSI-HAR, and the convolution-augmented Transformer within 24 kB and 54 kB, all far below the roughly 150 kB contiguous internal block that the runtime makes available with the radio stacks disabled. A full deep CNN and a Transformer therefore run on the bare classic ESP32: the weights (76 to 103 kB int8) live in the 4 MB flash, and only the activations occupy SRAM. The Chebyshev-KAN is the exception among the feed-forward models: it converts to int8, but `AllocateTensors` fails on the device even when the arena is grown to the full free block, so it does not deploy as built. This is a runtime limitation of its quantized polynomial and matmul layers rather than a raw memory ceiling.

The correction is therefore important: there is no general memory wall for feed-forward CSI models on the classic ESP32. The real barriers are the convertibility of the recurrent and state-space families and a runtime allocation failure for the KAN. What makes the tiny and classical tiers the practical choice is not that the deep models cannot fit, but that they are one to two orders of magnitude smaller and faster at comparable accuracy on the saturated benchmark (Section V-C).

C. ON-DEVICE MEASUREMENTS

Table 3 reports the measured latency and memory for the tiny and classical tiers, the models we recommend for deployment; the deep tier’s on-device outcome is in Table 2 (the CNN and Transformer allocate, the others do not). Every tiny network allocates with a peak tensor arena of only 7.7 to 13.4 kB, well under both the roughly 150 kB block the runtime exposes and even the deep CNN’s own 10 to 12 kB, and Figure 3 places all the deployable models against the budget. The tiny models are not chosen because the deep models fail to fit, but because they reach comparable accuracy on the saturated benchmark at a small fraction of the size and latency. Inference latency ranges from 14 ms for the 8-channel tiny CNN on CSI-HAR to 117 ms for the 32-channel tiny CNN on UT-HAR, all well within a real-time budget for activity recognition. A clear cost ordering appears: latency grows with channel width, and the 16-channel tiny CNN, the best accuracy-per-size point, costs 44 ms on UT-HAR and 30 ms on CSI-HAR. The tiny multilayer perceptron has the smallest arena of all because, despite its large weight file in flash, its activations are a single flat vector; this confirms that the runtime arena and not the model file size is what meets the memory budget. The classical learners, deployed as branch-only `emlearn C`, are the cheapest of all: the Decision Tree runs in about 1 microsecond and the Random Forest in 45 microseconds, roughly three orders of magnitude faster

TABLE 1. Deployable frontier: classical and tiny models on UT-HAR (fixed split, three seeds) and CSI-HAR (subject-independent leave-one-user-out, three folds). Accuracy is the mean $\pm$ standard deviation across runs. Complexity is tree node count, forest node count, or parameter count. The int8 size column applies to the neural models; trees and forests are deployed as emlearn C. All models in this table are small enough to be realistic candidates for the classic ESP32.

Dataset	Model	Tier	Accuracy (%)	Complexity	int8 size (kB)	Converts
UT-HAR	Decision Tree	classical	81.93 $\pm$ 0.09	481 nodes	emlearn C	n/a
	Random Forest	classical	91.93 $\pm$ 0.09	8720 nodes	emlearn C	n/a
	MLP (statistics)	classical	95.27 $\pm$ 0.52	14663 params	emlearn C	n/a
	Tiny MLP (net)	tiny-nn	87.60 $\pm$ 1.77	184.6 K params	184.0	yes
	Tiny CNN (8 ch)	tiny-nn	86.13 $\pm$ 1.95	5.8 K params	11.1	yes
	Tiny CNN (16 ch)	tiny-nn	94.93 $\pm$ 0.50	12.9 K params	18.7	yes
	Tiny CNN (32 ch)	tiny-nn	96.67 $\pm$ 0.50	31.0 K params	37.6	yes
CSI-HAR	Decision Tree	classical	41.43 $\pm$ 6.31	75 nodes	emlearn C	n/a
	Random Forest	classical	57.38 $\pm$ 3.21	1380 nodes	emlearn C	n/a
	MLP (statistics)	classical	58.10 $\pm$ 0.89	8583 params	emlearn C	n/a
	Tiny MLP (net)	tiny-nn	57.14 $\pm$ 2.92	106.8 K params	108.0	yes
	Tiny CNN (8 ch)	tiny-nn	54.52 $\pm$ 6.56	3.7 K params	9.1	yes
	Tiny CNN (16 ch)	tiny-nn	59.52 $\pm$ 2.36	8.7 K params	14.6	yes
	Tiny CNN (32 ch)	tiny-nn	62.14 $\pm$ 4.21	22.4 K params	29.3	yes

TABLE 2. Deep tier on both datasets with the measured on-device outcome. “Converts” is conversion to int8 TensorFlow Lite for Microcontrollers; “On-device” is the result of flashing each convertible model to the classic ESP32 (peak tensor arena from a successful `AllocateTensors`, or the failure mode). UT-HAR accuracy is from the companion single-pipeline benchmark; CSI-HAR accuracy is subject-independent (leave-one-user-out).

Model	Dataset	Accuracy (%)	int8 (kB)	Converts	On-device outcome
CNN	UT-HAR	96.87	103.0	yes	fits (10 kB arena)
BiGRU	UT-HAR	97.67	n/a	no	not convertible
Transformer	UT-HAR	98.27	88.2	yes	fits (24 kB arena)
Chebyshev-KAN	UT-HAR	95.67	90.4	yes	allocation fails
Lightweight-SSM	UT-HAR	90.20	n/a	no	not convertible
CNN	CSI-HAR	64.05	86.3	yes	fits (12 kB arena)
BiGRU	CSI-HAR	64.29	n/a	no	not convertible
Transformer	CSI-HAR	58.10	76.3	yes	fits (54 kB arena)
Chebyshev-KAN	CSI-HAR	65.71	73.7	yes	allocation fails
Lightweight-SSM	CSI-HAR	47.38	n/a	no	not convertible

than the tiny CNNs, with under 1 kB of working memory and no dynamic allocation. The energy column is an estimate (latency times the ESP32 active power; Section IV), not a direct measurement; even so, the tiny networks sit in the single-digit millijoule range per inference and the classical models are effectively free, which is the practically relevant comparison.

D. THE GENERALIZATION GAP: THE BINDING CONSTRAINT

The deployment results show that fitting a model on the device is largely solved. The harder question is whether a model works on a person it was not trained on. To isolate this, we evaluated the same models, features, and training pipeline on CSI-HAR under two protocols that differ only in how the data is split: a random stratified 80/20 split, in which recordings from every user appear in both train and test, and the subject-independent leave-one-user-out protocol used throughout this paper. Table 4 and Figure 7 report both.

The gap is large and consistent. Every model loses 25 to 33 percentage points when the test user is unseen, a mean drop of 29 points. The 32-channel tiny CNN reads 96.4% on the random split, which is in the range CSI-HAR papers commonly report, but only 63.1% leave-one-user-out; the

TABLE 3. Measured on-device results on the classic ESP32 (WiFi and Bluetooth disabled). Neural models run under TensorFlow Lite Micro; classical models run as emlearn C. Latency is the mean over 1000 invocations; peak RAM is the interpreter tensor arena for the neural models and the feature buffer for the classical models. Energy is *estimated* as latency times a representative active power of 0.165 W (50 mA at 3.3 V; the ESP32 datasheet [20] gives 30 to 68 mA at 240 MHz with the radio off), not directly measured. The deep CNN and Transformer also allocate on the device (Table 2); their per-inference latency was not profiled.

Model	Dataset	Latency (ms)	RAM (kB)	Energy (mJ, est.)
Decision Tree	UT-HAR	0.0011	<1	<0.001
Random Forest	UT-HAR	0.045	<1	0.007
Tiny CNN (8 ch)	UT-HAR	21.4	12.8	3.5
Tiny CNN (16 ch)	UT-HAR	44.4	13.0	7.3
Tiny CNN (32 ch)	UT-HAR	116.5	13.4	19.2
Tiny MLP (net)	UT-HAR	38.1	12.5	6.3
Tiny CNN (8 ch)	CSI-HAR	14.3	8.1	2.4
Tiny CNN (16 ch)	CSI-HAR	30.2	8.3	5.0
Tiny CNN (32 ch)	CSI-HAR	69.4	8.7	11.5
Tiny MLP (net)	CSI-HAR	22.5	7.7	3.7
Deep CNN/Transf.	both	fit, 10–54 kB arena (Table 2)		

Random Forest falls from 83.1% to 58.1%. Because nothing changes between the two columns except the split, the gap is not a property of any model or of its size. It is the cost of generalizing across people, and a random split hides it by letting




FIGURE 3. Measured peak tensor arena of the deployable models on the classic ESP32, against the roughly 150 kB contiguous block the runtime exposes (dashed line). Every convertible feed-forward model fits with a large margin, including the full deep CNN (10 to 12 kB) and Transformer (24 to 54 kB). The Chebyshev-KAN converts but fails to allocate, and the recurrent and state-space models do not convert; neither is shown. The tiny models are attractive for being far smaller and faster, not because the deep models fail to fit.




FIGURE 4. Accuracy versus int8 size for the tiny networks on both datasets (size on a logarithmic axis). The 16-channel tiny CNN sits at the top-left on both datasets, giving the best accuracy per kilobyte. The int8 file size is a lower bound on on-device cost; the runtime arena, which is what meets the device budget, is measured separately.




FIGURE 5. Accuracy of the classical and tiny tiers across UT-HAR and CSI-HAR. Accuracy is consistently lower on CSI-HAR, the single-ESP32 dataset with fewer samples per class.




FIGURE 6. Model complexity on a logarithmic axis (tree node count or network parameter count). The classical learners are one to three orders of magnitude smaller than the networks, which is why they fit the device with large margins.

TABLE 4. Cross-subject generalization gap on CSI-HAR. Identical models, features, and training; the only difference is the split. “Random” is a stratified 80/20 split over three seeds, in which each user appears in both train and test; “LOUO” is subject-independent leave-one-user-out over the three users. The gap is the accuracy lost when the test person is unseen.

Model	Random split (%)	LOUO (%)	Gap (pts)
Decision Tree	69.4	40.6	28.9
Random Forest	83.1	58.1	25.0
MLP (statistics)	85.2	56.9	28.3
Tiny CNN (16 ch)	92.1	62.1	29.9
Tiny CNN (32 ch)	96.4	63.1	33.3
Tiny MLP (net)	87.7	56.9	30.8
Mean	85.7	56.3	29.4

the model memorize each user’s channel signature. This is why we report subject-independent accuracy throughout, and it locates the open problem for CSI HAR on this hardware in generalization rather than deployment: the chip is no longer the limit.

VI. DISCUSSION

A. WHAT ACTUALLY RUNS ON A BARE ESP32

The frontier gives a direct answer to the title question, and it is more permissive than we first expected. On the PSRAM-




FIGURE 7. The cross-subject generalization gap on CSI-HAR. The same models score much higher on a random split (light bars) than under subject-independent leave-one-user-out (dark bars); the red number is the accuracy lost per model. The mean drop of 29 points is the cost of recognizing an unseen person, and it is invisible to a random split.

less classic ESP32 the deployable options for WiFi CSI HAR are broader than expected: they include the classical learners, the tiny convolutional networks, and a full deep CNN and convolution-augmented Transformer, which both allocate and run within tens of kilobytes. What the small models offer is efficiency, not exclusivity. On the saturated UT-HAR benchmark a classical multilayer perceptron on statistics reaches 95.3% and the 32-channel tiny CNN reaches 96.7% as an int8 model of 37.6 kB, within a point or two of the deep models but at a fraction of the size and latency. A practitioner should therefore start with the tiny and classical tiers for efficiency, knowing that a deep CNN or Transformer is also available if accuracy demands it; only the recurrent and state-space families and the Kolmogorov-Arnold model are off the table on this runtime. The 16-channel tiny CNN is a good default when latency matters (94.9% on UT-HAR at 44 ms and 13 kB of RAM), with the 32-channel model trading roughly double the latency for a small accuracy gain.

B. WHAT BLOCKS THE REST

The barriers are narrower than a blanket memory limit. Convertibility is the hard gate: recurrent and state-space models, attractive on the GPU, do not reach the microcontroller runtime at all without custom operators, because their fused recurrent and dynamic-tensor operators are outside the builtin set. The Chebyshev-KAN clears conversion but fails to allocate at runtime, a limitation of its quantized polynomial and matmul layers rather than a memory ceiling, since the much larger CNN and Transformer allocate comfortably on the same device. We initially assumed a memory wall blocked all deep models; flashing them to the board showed that assumption was wrong, which is the value of bare-device characterization over a GPU-only or PSRAM-equipped comparison. It also explains why prior on-device CSI work, navigating the same conversion limits, tends to favour convolutional models.

C. THE BINDING CONSTRAINT IS GENERALIZATION, NOT DEPLOYMENT

The two halves of the study point in the same direction. Once the deep CNN and Transformer are shown to run on the bare board, compute and memory are no longer what limits CSI HAR on this hardware. The leave-one-user-out results show what does: a mean 29-point drop when the test person is unseen (Section V-D). A practitioner who reads a 96% random-split number and provisions a tiny CNN will see about 63% on a new user, and no change of model or model size closes that gap, since the larger Transformer is no more robust across people than the tiny CNN. The practical consequence is that effort is better spent on generalization than on shrinking the model. A short per-user calibration or personalization pass, cheap given microsecond-scale classical inference and tens-of-milliseconds neural inference on the device, is the natural remedy, and cross-subject methods such as domain adaptation are the substantive open problem. The result also argues for reporting subject-independent accuracy

as standard practice, because a random split on a small multi-user corpus reports a number the deployed device will not reach.

D. WHY CLASSICAL MODELS DESERVE A SECOND LOOK

For the most constrained hardware, the smallest models are the natural design point rather than a compromise. The classical learners here are one to three orders of magnitude smaller than the networks, are exported to branch-only C with no runtime interpreter, run in microseconds on the device (about 1 microsecond for the Decision Tree and 45 microseconds for the Random Forest), and reach competitive accuracy on the saturated UT-HAR benchmark (95.3% for the statistics multilayer perceptron). For a sensing task that must share a microcontroller with an application, this predictability and small footprint are decisive.

E. LIMITATIONS

The frontier uses two datasets and a single microcontroller; the UT-HAR accuracy values inherit that benchmark's saturation, and the CSI-HAR numbers reflect a small three-user corpus evaluated subject-independently; here cross-subject generalization rather than deployability is the limiting factor. Separating the two would require larger multi-subject CSI datasets, which we leave to future work. A subject-independent accuracy in the 54 to 62% range is the expected difficulty of zero-shot cross-subject CSI sensing; in practice a short per-user calibration or personalization step, which is cheap to run on the device given the microsecond-scale classical inference, is the usual remedy and a natural extension of this work. The deep tier is reported on both datasets and its on-device outcome (convert, allocate, or fail) measured directly by flashing each model; the UT-HAR deep accuracies come from the companion single-pipeline benchmark. We report the deep models' on-device allocation and arena but not their per-inference latency, which we did not profile. The on-device latency and peak memory for the tiny and classical tiers are measured directly on the physical board; energy is estimated from the measured latency and the ESP32 datasheet active power rather than measured with a current probe, because the board carries no current-sense instrumentation, so the energy figures should be read as order-of-magnitude estimates.

VII. CONCLUSION

We presented a deployability frontier for WiFi CSI human activity recognition on the cheapest and most constrained commodity WiFi microcontroller, the PSRAM-less classic ESP32. Across two datasets and a graded set of models from classical learners through tiny networks to five deep architectures, we showed by flashing each to the board that the recurrent and state-space models do not convert to the microcontroller runtime and the Kolmogorov-Arnold model fails to allocate, but that a full deep CNN and a Transformer do run on the bare classic ESP32, alongside the tiny

and classical models. There is no general memory wall for feed-forward CSI models on this device. On the saturated UT-HAR benchmark the small models are competitive (a tiny CNN at 96.7% and a classical multilayer perceptron at 95.3%) at a fraction of the cost, while a subject-independent evaluation on CSI-HAR exposes the binding constraint: the same models drop by a mean of 29 percentage points (the best from 96% to 63%) when the test person is unseen, so cross-subject generalization, not deployability, is the open problem at this scale. We measured on the physical device that the tiny networks allocate with an 8 to 13 kB arena and run in 14 to 117 ms per inference, and that the deep CNN and Transformer allocate within 10 to 54 kB. The contribution is a practitioner-oriented, reproducible map of what is deployable on the bare device, with all code, models, and measurement firmware released. Future work includes a precise on-device energy study with a current probe and extending the frontier to additional microcontrollers, datasets, and live-capture validation.

ACKNOWLEDGEMENTS

The authors thank the Faculty of Information Technology, University of Central Punjab, for support and computing resources.

REPRODUCIBILITY

The training and quantization notebook, the exported int8 models and classical models, the emlearn C export of the trees, and the ESP32 measurement firmware and protocol will be released at publication. All off-device results are reproducible from the released notebook on standard cloud GPUs.

DATA AVAILABILITY

This study uses two public datasets: UT-HAR, available through the SenseFi library [3], and the CSI-HAR dataset [25]. No new human-subject data were collected. The code, trained models, and firmware that support the findings will be released in a public repository at publication.

AUTHOR CONTRIBUTIONS

A. W. Malik: conceptualization, methodology, supervision, project administration, and writing (review and editing). M. Ahmad: methodology, software, validation, formal analysis, investigation, data curation, writing (original draft), and visualization.

ETHICS DECLARATION

This work involves only secondary analysis of publicly available, de-identified datasets and on-device latency and memory measurements. No experiments on human participants or animals were conducted by the authors, so ethics approval was not required.

CONFLICT OF INTEREST

The authors declare no conflict of interest.

FUNDING

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

USE OF AI TOOLS

An AI assistant was used to help draft and edit code and prose. All experimental results, on-device measurements, and scientific claims were produced and verified by the authors, who take full responsibility for the content.

REFERENCES

- [1] D. Halperin, W. Hu, A. Sheth, and D. Wetherall, "Tool release: gathering 802.11n traces with channel state information," *ACM SIGCOMM Comput. Commun. Rev.*, vol. 41, no. 1, p. 53, Jan. 2011, doi: 10.1145/1925861.1925870.
- [2] Y. Ma, G. Zhou, and S. Wang, "WiFi sensing with channel state information: a survey," *ACM Comput. Surv.*, vol. 52, no. 3, pp. 1–36, Jun. 2019, doi: 10.1145/3310194.
- [3] J. Yang, X. Chen, H. Zou, D. Wang, Q. Xu, and L. Xie, "SenseFi: a library and benchmark on deep-learning-empowered WiFi human sensing," *Patterns*, vol. 4, no. 3, p. 100703, Mar. 2023, doi: 10.1016/j.patter.2023.100703.
- [4] Z. Chen, L. Zhang, C. Jiang, Z. Cao, and W. Cui, "WiFi CSI based passive human activity recognition using attention based BLSTM," *IEEE Trans. Mobile Comput.*, vol. 18, no. 11, pp. 2714–2724, Nov. 2019, doi: 10.1109/TMC.2018.2878233.
- [5] B. Li, W. Cui, W. Wang, L. Zhang, Z. Chen, and M. Wu, "Two-stream convolution augmented transformer for human activity recognition," in *Proc. AAAI Conf. Artif. Intell.*, vol. 35, no. 1, 2021, pp. 286–293, doi: 10.1609/aaai.v35i1.16103.
- [6] M. S. Khan et al., "Human activity recognition through augmented WiFi CSI signals by lightweight attention-GRU," *Sensors*, vol. 25, no. 5, p. 1547, Mar. 2025, doi: 10.3390/s25051547.
- [7] T. D. Quy, C.-Y. Lin, and T. K. Shih, "Enhanced human activity recognition using Wi-Fi sensing: leveraging phase and amplitude with attention mechanisms," *Sensors*, vol. 25, no. 4, p. 1038, Feb. 2025, doi: 10.3390/s25041038.
- [8] Z. Liu et al., "KAN: Kolmogorov-Arnold networks," *arXiv:2404.19756*, 2024.
- [9] X. Xu, J. Yang et al., "Neurosymbolic AI empowered consumer electronics healthcare for WiFi-based human activity recognition," *IEEE Trans. Consum. Electron.*, vol. 71, no. 4, pp. 12056–12067, Nov. 2025.
- [10] M. Alikhani, "KAN-HAR: a human activity recognition based on Kolmogorov-Arnold network," *arXiv:2508.11186*, 2025.
- [11] A. Gu and T. Dao, "Mamba: linear-time sequence modeling with selective state spaces," *arXiv:2312.00752*, Dec. 2023.
- [12] J. Liu, Y. Huang, X. Shi, X. Ren, T. Mi, and R. C. Qiu, "TF-Mamba: a lightweight state-space model for Wi-Fi-based human activity recognition," *IEEE Sensors J.*, 2025, IEEE Xplore document 10817504.
- [13] H. Tan, Y. Dao, H. Zhang, T. Guo, and W. Wang, "WiFi signal-based human activity recognition using bidirectional Mamba models," in *Proc. IEEE 11th World Forum Internet Things (WF-IoT)*, 2025, doi: 10.1109/WF-IoT64238.2025.11270783.
- [14] H. Xu et al., "MambaLite-Micro: memory-optimized Mamba inference on MCUs," *arXiv:2509.05488*, 2025.
- [15] S. M. Hernandez et al., "Tools and methods for achieving Wi-Fi sensing in embedded devices," *Sensors*, vol. 25, no. 19, p. 6220, Oct. 2025, doi: 10.3390/s25196220.
- [16] K. Liu, "STAR: a privacy-preserving, energy-efficient edge AI framework for human activity recognition via Wi-Fi CSI in mobile and pervasive computing environments," *arXiv:2510.26148*, 2025.
- [17] Z. Wen, Y. Ruan, X. Wang, and J. Zhou, "ESP-Fi HAR: a low-power WiFi CSI dataset for ad-hoc IoT human activity recognition," *Pervasive Mob. Comput.*, 2026, doi: 10.1016/j.pmcj.2026.102058.
- [18] R. Koo, X. Liang, D. Mishra, and A. Seneviratne, "A survey on green wireless sensing: energy-efficient sensing via WiFi CSI and lightweight learning," *Energies*, vol. 19, no. 2, p. 573, 2026, doi: 10.3390/en19020573.
- [19] T. Suwannaphong et al., "Optimising TinyML with quantization and distillation of transformer and Mamba models for indoor localisation on edge devices," *Sci. Rep.*, 2025, *arXiv:2412.09289*.

- [20] Espressif Systems, “ESP32 series datasheet,” v5.2, 2025. [Online]. Available: https://www.espressif.com/sites/default/files/documentation/esp32_datasheet_en.pdf
- [21] J. Nordby et al., “emlearn: machine learning inference engine for microcontrollers and embedded devices,” software, 2019, doi: 10.5281/zenodo.2589394.
- [22] M. Ali et al., “A comparison of machine learning algorithms for Wi-Fi sensing using CSI data,” *Electronics*, vol. 12, no. 18, p. 3935, 2023, doi: 10.3390/electronics12183935.
- [23] B. Saha, R. Samanta, S. K. Ghosh, and R. B. Roy, “TinyTNAS: GPU-free, time-bound, hardware-aware neural architecture search for TinyML time series classification,” *arXiv:2408.16535*, 2024.
- [24] T. King, Y. Zhou, T. Röddiger, and M. Beigl, “MicroNAS for memory and latency constrained hardware aware neural architecture search in time series classification on microcontrollers,” *Sci. Rep.*, vol. 15, 2025, doi: 10.1038/s41598-025-90764-z.
- [25] S. Ghorai, “CSI-HAR dataset: WiFi channel state information for human activity recognition,” *Kaggle dataset*, 2024. [Online]. Available: <https://www.kaggle.com/datasets/sayakghorai34/csi-har-dataset>
- [26] TensorFlow project, “Conversion of RNN/LSTM to TensorFlow Lite requires Select-TF-ops,” *GitHub issue 36688*, 2020.
- [27] TensorFlow Lite Micro project, “Unsupported dynamic-tensor and TensorList operators in the microcontroller runtime,” *GitHub issues 907 and 1570*, 2021.

